

EXP NO: 4b

Date: 20-8-26

ENSEMBLE METHODS: RANDOM FOREST

AIM:

To implement a **Random Forest classifier** for a binary classification task, tune important hyperparameters, evaluate its performance using classification metrics, and identify the most important features.

ALGORITHM:

1. Import the required Python libraries.
2. Load the **Breast Cancer Wisconsin dataset**.
3. Split the dataset into **80% training and 20% testing** data using stratification.
4. Use the raw features since Random Forest does not require feature scaling.
5. Create a RandomForestClassifier.
6. Perform hyperparameter tuning using GridSearchCV.
7. Train the Random Forest model using the best parameters.
8. Predict the test data and calculate **accuracy, precision, recall, F1-score, ROC-AUC**, and confusion matrix.
9. Calculate feature importance and display the **top 10 important features**.
10. Plot the ROC curve and feature importance.

PROGRAM:

```
# 1) Import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix,
    classification_report,
    roc_auc_score,
    roc_curve
)

# 2) Load dataset
data = load_breast_cancer()

X = pd.DataFrame(
    data.data,
    columns=data.feature_names
)

y = pd.Series(
    data.target,
    name="target"
)
```

```

# 3) Train/Test split
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.20,
    random_state=42,
    stratify=y
)

# 4) Define Random Forest model
rf = RandomForestClassifier(
    random_state=42,
    n_jobs=-1
)

# 5) Hyperparameter tuning
param_grid = {
    "n_estimators": [100],
    "max_depth": [None, 10],
    "min_samples_split": [2],
    "min_samples_leaf": [1]
}

grid = GridSearchCV(
    estimator=rf,
    param_grid=param_grid,
    scoring="f1",
    cv=3,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)

print("Best Parameters (CV):")
print(grid.best_params_)

best_rf = grid.best_estimator_

# 6) Train final model and predict
best_rf.fit(X_train, y_train)

y_pred = best_rf.predict(X_test)
y_prob = best_rf.predict_proba(X_test)[:, 1]

# 7) Evaluation
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, zero_division=0)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)
cm = confusion_matrix(y_test, y_pred)

```

```

print("\n=== Random Forest - Test Metrics ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall   : {rec:.4f}")
print(f"F1-Score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f}")

print("\nConfusion Matrix:")
print(cm)

print("\nClassification Report:")
print(
    classification_report(
        y_test,
        y_pred,
        zero_division=0
    )
)

# 8) Feature Importance - Top 10
importances = pd.Series(
    best_rf.feature_importances_,
    index=X.columns
)

top10 = importances.sort_values(
    ascending=False
).head(10)

print("\nTop 10 Important Features:")
print(top10)

# Plot feature importance
plt.figure(figsize=(10, 6))
top10.sort_values().plot(kind="barh")

plt.xlabel("Importance")
plt.ylabel("Feature")
plt.title("Top 10 Feature Importances - Random Forest")
plt.grid(True)
plt.show()

# 9) ROC Curve
fpr, tpr, thresholds = roc_curve(
    y_test,
    y_prob
)

```

```
plt.figure()

plt.plot(
    fpr,
    tpr,
    label=f"Random Forest (AUC = {auc:.3f})"
)

plt.plot(
    [0, 1],
    [0, 1],
    linestyle="--",
    color="gray"
)

plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve – Random Forest")

plt.legend()
plt.grid(True)
plt.show()
```

OUTPUT:

```
Best Parameters (CV):
{'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
```

```
=== Random Forest – Test Metrics ===
```

```
Accuracy : 0.9561
Precision: 0.9589
Recall   : 0.9722
F1-Score : 0.9655
ROC-AUC  : 0.9937
```

```
Confusion Matrix:
```

```
[[39  3]
 [ 2 70]]
```



```
Classification Report:
```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	42
1	0.96	0.97	0.97	72
accuracy			0.96	114
macro avg	0.96	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

```
Top 10 Important Features:
```

```
worst area          0.140016
worst concave points 0.129530
worst radius        0.097696
mean concave points  0.090885
worst perimeter     0.072226
mean perimeter      0.069574
mean radius         0.068676
mean concavity      0.057638
mean area           0.049172
worst concavity     0.034340
dtype: float64
```

Output: 1.000000



RESULT:

The **Random Forest classifier** was successfully implemented and optimized using GridSearchCV. The model achieved high classification performance based on accuracy, precision, recall, F1-score, and ROC-AUC. The **confusion matrix**, **ROC curve**, and **top 10 important features** were also obtained successfully.